

Wrought Iron Commands

To compete with enterprise tools (or Python scripts written by senior engineers), every single command must expose the full power of the underlying library (`pandas`, `scikit-learn`, `cryptography`, etc.) through CLI flags.

We call this the **"Surface-to-Core" API philosophy**: The default usage is simple (`wi clean ml-impute users birth_year`), but the *optional* flags expose the full Core API of the Python library.

Here are 5 concrete examples showing this "Deep Argument" structure applied to different modules, proving that this depth extends far beyond just machine learning.

1. The Geo Module: Deep Clustering

Instead of just "grouping points," we expose the full **DBSCAN** algorithm parameters from `scikit-learn`.

```
wi geo cluster users \
--eps 0.5 \                      # (DBSCAN) Max distance between points (km)
--min-samples 10 \                  # (DBSCAN) Min points to form a cluster
--metric haversine \                # (DBSCAN) Use earth-curvature math, not Euclidean
--algorithm ball_tree \             # (Optimization) Faster spatial indexing method
--n-jobs -1 \                      # (Performance) Use all CPU cores
--out-col catchment_zone \          # (Save) Name of the new column
--noise-label "Rural" \             # (Labeling) What to call unclustered points
--dry-run                            # Preview cluster sizes without saving
```

2. The Aggregate Module: Deep Statistics

Instead of just "grouping," we expose the full power of Pandas `groupby()` and `agg()`.

```
wi aggregate groupby transactions amount \
--by suburb,category \            # Multi-level grouping
--agg mean,sum,std,max \          # Multiple statistical operations
--sort-by sum \                   # Order result by total amount
--ascending False \               # Highest first
--min-count 5 \                  # Drop groups with <5 records (Privacy)
--dropna \                        # Exclude groups where 'suburb' is NULL
--format markdown \               # Output format (Table, CSV, Markdown, JSON)
--out report_Q3.md                # Save directly to file
```

3. The Query Module: Deep Filtering

Exposing the granularity of Pandas querying and string handling.

```
wi query search students "Smith" \
--col name,parent_name \ # Limit search scope
--case-sensitive False \ # Ignore capitalization
--regex \ # Treat "Smith" as Regex pattern
--fuzzy-threshold 85 \ # Allow slight misspellings ("Smyth")
--context 50 \ # Show 50 chars of text around the match
--limit 100 \ # Stop after 100 hits
--output-mode id-only # Only return IDs for piping to another command
```

4. The Ops Module: Deep Scheduling

Exposing system-level cron/task controls.

```
wi ops schedule create "clean-nightly" \
--cmd "wi clean impute-group users" \
--cron "0 3 * * *" \ # Run at 3:00 AM daily
--timeout 3600 \ # Kill if runs longer than 1 hour
--retry 3 \ # Retry 3 times if it fails
--on-fail-email "admin@secure.mil" \ # Local alert hook
--cpu-limit 50% \ # Nice value (don't hog the server)
--log-level DEBUG # Verbose logging for debugging
```

5. The Audit Module: Deep Forensics

Exposing the cryptographic controls.

```
wi audit hash-verify users \
--algorithm sha256 \ # Hashing standard
--exclude-cols updated_at \# Ignore dynamic timestamp columns
--salt "random-seed" \ # Salt the hash for extra security
--strict \ # Fail if schema changed (not just data)
--chunk-size 10000 \ # Optimize memory for massive tables
--report-format pdf \ # Generate formal certificate
--signer-key private.pem # Digitally sign the verification report
```

The "Help" Experience

Because these commands are deep, the Developer Experience (DX) must be perfect. Using `Typer`, typing `--help` for any command will auto-generate this documentation in the terminal:

```
$ wi clean ml-impute --help

Usage: wi clean ml-impute [OPTIONS] TABLE COL

    Impute missing values using K-Nearest Neighbors (KNN).

Arguments:

    TABLE  Target database table  [required]
    COL    Column with missing values  [required]

Options:

    --method [knn|iterative]  Imputation strategy [default: knn]
    --neighbors INT          Number of neighbors (k) [default: 5]
    --weights [uniform|distance]
                                Weight function used in prediction [default: uniform]
    ...
...
```

This confirms that **Wrought Iron** is not just a wrapper; it is a full-featured interface to the Python Data Science stack.